

Standardizing RNA-seq Analysis of Fungal Pathogens Using BRC-Analytics and Agentic AI: A *Candidozyma auris* Case Study

Anton Nekrutenko¹, Danielle Callan², Marius Van Den Beek¹, Dannon Baker³, David Rogers⁴, Aysam Guerler³, John Chilton¹, Hiram Clawson⁵, Scott Cain¹, Teresa O'Meara⁶, Kelsey Beavers⁷, Jeremy Goecks⁹, Michael Schatz³, Maximilian Haeussler⁵, Bjorn Gruning⁸, and Sergei Kosakovsky Pond²

¹ Dept. of Biochemistry and Molecular Biology, The Pennsylvania State University, University Park, PA, USA

² Dept. of Biology, Temple University, Philadelphia, PA, USA

³ Dept. of Biology, Johns Hopkins University, Baltimore, MD, USA

⁴ Clever Canary, LLC, Santa Cruz, CA, USA

⁵ Baskin School of Engineering, University of California, Santa Cruz, USA

⁶ Dept. of Microbiology and Immunology, University of Michigan, Ann Arbor, MI, USA

⁷ Texas Advanced Computing Center, The University of Texas, Austin, TX, USA

⁸ Dept. of Bioinformatics, Albert-Ludwigs-University Freiburg, Freiburg, Baden-Württemberg, Germany

⁹ Moffitt Cancer Center, Tampa, FL, USA

Correspondence should be addressed to AN and SKP: aun1@psu.edu, spond@temple.edu

Abstract

Candidoizyia auris (*C. auris*) has emerged as a critical global health threat due to multidrug resistance and healthcare-associated transmission. While RNA-seq has become the primary tool for studying *C. auris* pathogenesis, inconsistent use of reference genomes and bioinformatics tools complicate cross-study comparisons. Here we demonstrate how BRC-Analytics, a platform for pathogen genomics, combined with an agentic AI assistant (Claude Code Agent), enables reproducible RNA-seq analysis. By re-analyzing data from two publications we achieved near-perfect correlation with published results despite annotation version differences. We addressed provenance challenges associated with using AI agents with Galaxy by forcing them to invoke Galaxy’s native tools rather than manipulating data directly. For custom analyses outside Galaxy’s toolset, we provide standalone JupyterLite notebooks that reproduce our analysis without AI involvement. This framework—combining AI-assisted automation with rigorous provenance tracking—establishes a template for standardized, reproducible fungal pathogen genomics. To the best of our knowledge, this is the first example of integration between public data repositories, reproducible analysis workflows, and agentic AI tools. Our subsequent efforts will focus on improving the seamlessness of this integration.

Introduction

Candidoizyia auris (formerly *Candida auris*; NCBI:txid498019) represents one of the most urgent antimicrobial resistance threats facing global health systems. First isolated from external ear canal of Japanese hospital patient in 2009 [1], this fungal pathogen has since spread worldwide. CDC classifies *C. auris* as an urgent threat—the first fungal pathogen to receive this designation—due to multidrug resistance (often to all major antifungal classes), healthcare-associated transmission, and 30-60% mortality rates [2,3]. *C. auris* persists on surfaces, colonizes skin, and forms biofilms on medical devices, enabling difficult-to-control nosocomial outbreaks [3]. WHO designates *C. auris* as critical-priority fungal pathogen [4], and NIAID has prioritized development of new therapeutics [5].

Compared to other key human pathogens (such as a SARS-CoV-2 or HIV, for example) the amount of publicly available sequence data for *C. auris* is modest (Table 1). Two categories of projects account for 98% of all data: whole genome sequencing efforts (WGS) and RNA-seq projects. The WGS data are mostly derived from outbreak surveillance efforts conducted by various state public health agencies (Supp. Table 1). The majority of RNA-seq data on the other hand are produced by academic research labs. This reflects the importance of transcriptomic analyses to understanding the fundamental biology of this pathogen. While whole-genome sequencing dominates by run count (26,201 WGS vs 812 RNA-seq runs; 96.3% vs 3.0%), 64 of 237 *C. auris* BioProjects (27%) are RNA-seq studies. This disparity reflects study design: WGS projects sequence many isolates for outbreak surveillance (average 156 runs/project), whereas RNA-seq examines specific biological conditions (average 13 runs/project). Given RNA-seq accounts for over one-quarter of *C. auris* research projects, standardizing analysis is a critical priority.

Table 1: Summary of *C. auris* sequencing data in NCBI SRA (December 2025). BioProject is an NCBI database entry grouping related sequencing runs from a single study. Assay types: WGS = whole genome sequencing; RNA-Seq = transcriptome sequencing; AMPLICON = targeted amplicon sequencing; WGA = whole genome amplification; miRNA-Seq = microRNA sequencing; ChIP-Seq = chromatin immunoprecipitation sequencing; Tn-Seq = transposon insertion sequencing; Targeted-Capture = hybridization capture sequencing; WCS = whole chromosome sequencing; Bisulfite-Seq = DNA methylation sequencing.

Assay Type	BioProjects	Total Runs	Total Bases	Avg Runs/Project
WGS	168	26,201	45.6 Tb	156.0
RNA-Seq	64	812	4.7 Tb	12.7
AMPLICON	4	87	17.2 Gb	21.8
WGA	2	38	34.3 Gb	19.0
miRNA-Seq	1	24	4.7 Gb	24.0
ChIP-Seq	2	14	112.5 Gb	7.0
OTHER	2	13	40.0 Gb	6.5
Tn-Seq	1	6	19.7 Gb	6.0
Targeted-Capture	1	5	2.8 Gb	5.0
WCS	1	1	2.1 Gb	1.0
Bisulfite-Seq	1	1	383.6 Mb	1.0
TOTAL	237	27,202	50.5 Tb	

To understand the analytical landscape of *C. auris* transcriptomic studies we surveyed all available RNA-seq data associated with that species. Specifically, for all 64 RNA-seq BioProjects listed in Table 1 we attempted to retrieve associated publications. Of 64 BioProjects, 20 (31%) had linked manuscripts (21 papers total, 2018-2025) while 44 remained unpublished or in pre-print stage. For papers with available full text (17/20), we extracted reference genome and analysis tool information (Table 2; also see Supp. Table 2).

Table 2: RNA-seq methodology across 20 published *C. auris* studies with linked BioProjects

Category	Finding
Reference Genome	B8441/GCA_002759435.x (12/20, 60%); multiple clades (5/20); not specified (2/20)
Alignment Tool	HISAT2 (7), STAR (5), Bowtie2 (4), BWA (3), TopHat2 (1)
Quantification	featureCounts (5), HTSeq (4), StringTie (2), Kallisto (2), RSEM (1)
DE Analysis	DESeq2 (12), edgeR (4), Cufflinks (1)
Publication Years	2018 (2), 2021 (4), 2022 (4), 2023 (2), 2024 (5), 2025 (4)

Despite tool convergence, reference genome usage remains inconsistent. While 60% of published studies use B8441 (GCA_002759435 family), annotation versions vary—some cite only “B8441” without version, others specify GCA_002759435.2 or GCA_002759435.3. This creates reproducibility challenges (e.g., gene identifiers differ between versions) and complicates interpretation of published data in context of new genomes and vice versa. Similarly, tool version reporting is frequently incomplete or absent—papers cite “HISAT2” or “DESeq2” without specifying version numbers, yet algorithm behavior and output can differ substantially between releases. Without precise version information, reproducing published results becomes guesswork, undermining scientific rigor. These findings underscore need for standardized platforms specifying precise genome versions, tool versions, and parameters.

Here, we demonstrate how a new environment for the analysis of pathogen, host, and vector data—BRC-Analytics (<https://brc-analytics.org>)—can be used for standardizing and simplifying RNA-seq analyses using two recent *C. auris* studies as an example. Our approach makes cutting edge tools and powerful computational infrastructure freely accessible to any biologist. Importantly, the combination of BRC-analytics, the Galaxy platform, and Agentic AI tools built on Large Language Models (LLM) tools described here automatically keeps provenance and ensures analytical reproducibility: any analysis conducted within our system can be understood and replicated by others.

Results

BRC-Analytics

BRC-Analytics (<https://brc-analytics.org>) is an online, browser-based analysis environment designed to make comprehensive and reproducible genomic analyses of infectious diseases accessible to everyone. Developed under the NIAID-funded Bioinformatics Resource Centers (BRCs) program, it leverages the Galaxy workflow system to enable users to begin with raw sequencing reads and achieve publication-ready results without the need for local software installations or manual data transfers between tools. The platform integrates authoritative genomic data from multiple sources: NCBI Datasets provides reference genomes (currently 1,601 assemblies for 1,149 pathogen, host, and vector taxa, with continuous expansion planned), UCSC Genome Browser supplies genome annotations including gene coordinates and regulatory elements, and EBI ENA facilitates access to public sequence read archive data through local caching for quick searches. BRC-Analytics pairs these data sources with community-curated best-practice analysis workflows covering essential steps like quality control, read mapping, variant identification, and annotation. Galaxy serves not only for launching and running workflows but also as an environment for interpretive analyses through interactive tools like Jupyter, RStudio, and JupyterLite. The platform utilizes cloud-based computation, versioned workflows, and interactive visualizations to create a seamless, reproducible interface. The substantial computational and storage resources required are provided by ACCESS-CI infrastructure in the US, with BRC-Analytics and Galaxy hosted on servers at the Texas Advanced Computing Center (TACC). This approach unifies data and analytical capabilities, making advanced pathogen genomics available to a wider research community.

Two representative studies

The *Introduction* section above described a survey of all publicly available *C. auris* sequence data with a particular focus on RNA-seq studies and associated publications (Supp. Table 2). From these publications we selected two studies. The first, Santana et al. (2023), identified *SCF1* gene as *C. auris*-specific adhesin essential for biofilm formation and virulence (PRJNA904261) [6]. The second, Wang et al. (2024), showed that glycan-lectin interactions modulate colonization and fungemia (PRJNA1086003) [7]. These two studies are good representatives of *C. auris* RNA-seq methodology. Both use B8441 (Clade I) reference genome, which dominates the field (14/20 published studies). Wang employs HISAT2/STAR + DESeq2, the most common pipeline (DESeq2 in 13/20, HISAT2 in 6/20 studies). Sample sizes of 13 and 6 runs bracket the typical range (median ~13-15, with 5 studies having exactly 6 runs). As 2023-2024 publications, they reflect current practices unlike older studies using outdated tools (TopHat2, Cufflinks). Both study adhesion/biofilm phenotypes, the dominant research theme alongside drug resistance.

Configuring analysis

The workflow invoked in the previous section created read counts for each gene annotated in *C. auris* assembly (GCA_002759435.3) for all samples described in each of the papers (Supp. File 1). In order to continue with the analysis of differential gene expression we now need to organize data in way that reflects the experimental design of the two manuscripts (To make this section more understandable for the reader this section focuses on the re-analysis of Santana et al. [6] data. Data from Wang et al. (2024) data [7]) was re-analyzed in the exactly the same way). This implies identifying which samples belong to which condition, identifying replicates and so on. While this can be done manually we can harness the power of large language models (LLMs) to simplify this step. Before proceeding to description of our computational set-up we note that any results produced by LLMs must always be verified. This is why for each analyses we always as an LLM agent to first produce a plan of action. We then review this plan and modify as necessary before allowing the agent to proceed.

For this analyses we used Claude Code Agent (CCA) produced by Anthropic. It was configured to interact with Galaxy platform via application programming interface (API) key that effectively allows CCA to take actions on behalf of the user. CCA was running on a local computer we used to prepare this manuscript and interacted with Galaxy via Internet. The precise record of our interaction with CCA and its interaction with Galaxy is given in Supp. File 2 (also see Video 2 and Materials and Methods for the description of our setup).

Organizing data

Authors of each paper have deposited sequencing data into NCBI SRA and were given a BioProject identifier—an entity grouping related sequencing runs from a single study. Before performing the differential expression analyses we need to understand how samples deposited to SRA correspond to experimental conditions described in each manuscript. To begin we deposited PDFs of each manuscript along with all available supplemental data into a separate folder on a local computer. We then provided the following instructions to the CCA:

I need to split Galaxy dataset collection #244 into several collections corresponding to experimental conditions described in the manuscript (check pdfs and xlsx files in this directory). In order to do this you need to download metadata for sequencing runs for bioproject PRJNA904261 to obtain accessions and metadata. You should then figure out how SRA accessions correspond to experimental conditions described in the paper. You should then present these finding to me, so that I can tell you what to do next.

In the above prompt we specifically mentioned “Collection #244”—a Galaxy artifact containing read counts for all samples described in this study (can be viewed at <https://usegalaxy.org/u/cartman/h/prjna904261-perm>). A plan generated in response to this prompt can be seen in Supp. File 2. One can see that CCA correctly identified the relationship between datasets and experimental conditions described in the manuscript.

After reviewing the plan we instructed CCA to enact it:

Go ahead and execute the plan. Once you are done please add name tags to dataset collection containing data we need to used for DeSeq2 analysis. E.g., label collections with names tags such as AR0382_WT, AR0387_WT, and tnSWI1.

This step generated three dataset collections in Galaxy history corresponding to the three conditions

described in the paper: AR0382_WT, AR0387_WT, and tnSWI1.

We then repeated this procedure in a separate Galaxy history containing read count derived from Wang et al. 2024 [7].

Table 3: Breakdown of datasets for DESeq2 analysis. For Santana et al. AR0382_WT/tnSWI1 and AR0382_WT/AR0387_WT comparisons were performed. For Wang et al. AR0382 in vitro/AR0387 in vitro and AR0382 in vivo/AR0387 in vivo comparisons were performed.

Study	Condition	SRR Accessions	Description
Santana et al.	AR0382_WT	SRR22376031, SRR22376032	Wild-type reference
	AR0387_WT	SRR22376029, SRR22376030	Poorly adhesive strain
	tnSWI1	SRR22376027, SRR22376028	SWI1 mutant
Wang et al.	AR0382 in vitro	SRR28790270, SRR28790272, SRR28790274	In vitro culture
	AR0387 in vitro	SRR28790276, SRR28790278, SRR28790280	In vitro culture
	AR0382 in vivo	SRR28791430, SRR28791431, SRR28791432	In vivo infection
	AR0387 in vivo	SRR28791433, SRR28791434, SRR28791437, SRR28791438	In vivo infection

Expression analysis and interpretation

In the previous section we have configured our data so we can perform differential expression. We then re-ran differential expression with DESeq2 on data from both manuscripts as described in Table 3 and performed a systematic comparison of log2 fold changes against published supplementary data.

A technical challenge arose from differences in genome annotation versions. Both published studies used an older *C. auris* annotation with 6-digit gene ID suffixes (e.g., B9J08_001458), while we relied on the latest assembly (GCA_002759435.3) that uses 5-digit suffixes (e.g., B9J08_03708). Our re-analyses identified comparable DEG counts to the published results. Yet because different genome versions were used we could not map gene names directly. We decided to perform log2 fold-change (LFC) correlation mapping: for each published differentially expressed gene (DEG) in the two papers, we matched genes in our data by finding the closest log2 fold-change value in our results, accepting matches within 0.1 LFC units. The method achieved >95% mapping success with mean differences of only 0.01-0.02. CCA performed this entire procedure using the following prompt (here, again, using Santana et al. as an example) :

Datasets #521 and #523 in <https://usegalaxy.org/u/cartman/h/prjna904261-perm> represent DESeq2 results for AR0382_WT/tnSWI1 and AR0382_WT/AR0387_WT comparisons, respectively. Compare them with the results reported in Santana et al. using paper PDF and supplementary files. Resolve gene name mapping using log2 fold-change (LFC) correlation analysis

Comparison with Santana et al. (2023) results

The study compared three strains (Table 3): AR0382_WT, a wild-type highly adhesive Clade I isolate; AR0387_WT, a poorly adhesive clinical isolate; and tnSWI1, a transposon-insertion mutant

of AR0382 with disrupted SWI1, a chromatin remodeling factor. Each strain was sequenced in duplicate. The first comparison (Fig 1D) examined the tnSWI1 mutant versus wild-type AR0382 to identify genes affected by SWI1 disruption. The second comparison (Fig S5A) contrasted AR0387 against AR0382 to characterize expression differences between adhesive and non-adhesive strains. Both comparisons yielded excellent validation metrics. For the first (tnSWI1/AR0382_WT) comparison, we successfully mapped 1,186 differentially expressed genes and obtained a Pearson correlation of $r = 0.9971$ ($R^2 = 0.9953$) with 99.7% direction agreement. The second comparison (AR0382_WT/AR0387_WT) mapped 1,557 genes with $r = 0.9883$ ($R^2 = 0.9768$) and perfect 100% direction agreement. SCF1, the central finding of the Santana study, was the most strongly down-regulated gene in both comparisons. The published analysis reported SCF1 (B9J08_001458) with log2 fold changes of -6.68 (Fig 1D) and -7.25 (Fig S5A). Our reanalysis identified the corresponding gene (B9J08_03708) with log2 fold changes of -6.82 and -7.35, respectively, confirming the paper’s key finding with minimal deviation.




Figure 1: Validation of Santana et al. DESeq2 results. Left: tnSWI1 vs AR0382 comparison (Fig 1D). Right: AR0387 vs AR0382 comparison (Fig S5A). Red dashed line indicates perfect correlation ($y=x$).

Comparison with Wang et al. (2024) results

This study compared two strains with distinct aggregation phenotypes: AR0382 (B11109), a highly aggregative biofilm-forming strain, and AR0387 (B8441), a non-aggregative strain. RNA-seq was performed under two conditions: in vitro biofilm growth (3 replicates per strain) and in vivo colonization of mouse jugular vein catheters (3 replicates for AR0382, 4 for AR0387). The authors reported 76 differentially expressed genes (DEGs) in the in vitro comparison and 259 DEGs in the in vivo comparison, using thresholds of $FDR < 0.01$ and $|LFC| \geq 1.0$. Our reanalysis achieved perfect correlation with the published results. For the in vitro condition, we identified 73 DEGs with Pearson $r = 1.0000$ and 100% direction agreement across 75 matched genes. The in vivo analysis reproduced all 259 DEGs with identical correlation metrics. The key adhesin genes highlighted in the paper showed excellent concordance. SCF1 exhibited LFC of 8.61 (paper) versus 8.67 (our analysis) in vitro, and 4.47 versus 4.53 in vivo. ALS4112 showed similarly close agreement: 5.07 versus 5.08 in vitro, and 2.56 versus 2.56 in vivo. Other adhesin family members (IFF4109), virulence-associated genes (SAP7), and drug resistance genes (MDR1) all matched within 0.01 LFC units.




Figure 2: Validation of Wang et al. DESeq2 results. Left: In vitro biofilm comparison ($n=75$ genes). Right: In vivo mouse catheter model ($n=259$ genes). Red dashed line indicates perfect correlation ($y=x$). Key adhesin genes SCF1 and ALS4112 are labeled.

Maintaining provenance

Integrating AI agents with analytical platforms like Galaxy presents a provenance challenge. When analysis alternates between Galaxy workflows and external AI-generated scripts, the chain of reproducibility breaks—Galaxy cannot track code executed outside its environment, and AI agents generate numerous artifacts (Python scripts, intermediate files) that are difficult to document systematically. To preserve provenance, we configured our AI agent to invoke Galaxy’s native tools through the API rather than manipulating data directly. When a CCA interacts with Galaxy via API, it can directly manipulate datasets and collections—but this bypasses Galaxy’s tool framework, losing reproducibility and workflow compatibility. By constraining the agent to use Galaxy tools (e.g., Apply Rules, Filter, DESeq2), all operations remain tracked in Galaxy histories, can be extracted into reusable workflows, and produce identical results on re-execution. However, some operations—such as the LFC-based gene mapping described above—require custom code that Galaxy does not natively support. To address this gap, we developed a standalone JupyterLite notebook (with the help of the same CCA) that reproduces the validation figures shown here without any AI involvement (notebooks can be accessed in Galaxy histories [8,9]). The notebook requires two inputs: (1) DESeq2 output from Galaxy (TSV with Gene_ID, log2FoldChange, padj columns) and (2) publication data reformatted as CSV with gene_id and log2fc columns. In our workflow, the AI agent’s role was limited to extracting and reformatting data from Excel supplementary files into this simple CSV structure—a step that becomes unnecessary when published supplementary data already conforms to standard formats. Ideally, the CCA functionality should be tightly integrated with BRC-Analytics and Galaxy—a current development priority for us (see below).

Discussion

Toward standardization of fungal genomics

Both studies achieved excellent validation status, with Santana et al. showing $R^2 > 0.97$ across both comparisons and Wang et al. achieving perfect correlation ($r = 1.00$) in both conditions. These results demonstrate that Galaxy-based reanalysis using standard workflows produces results

equivalent to published analyses, and that differential expression patterns are highly reproducible when using the same statistical methods and significance thresholds. From a technical perspective, this validation demonstrates that different genome annotation versions can be effectively mapped using LFC correlation, enabling retrospective validation of studies that used older reference genomes. Both studies employed consistent statistical thresholds ($\text{FDR} < 0.01$, $\text{LFC} \geq 1.0$), facilitating direct comparison of results.

Importance of LLMs and their responsible use

For reproducibility, LLM-assisted analyses should be conducted through agentic coding tools such as Claude Code or Gemini Code Assist rather than chat-based interfaces. These tools automatically track all generated artifacts—scripts, intermediate files, and analysis outputs—within version-controlled repositories (e.g., GitHub), creating a complete audit trail. While this workflow may currently seem complex for bench biologists unfamiliar with command-line interfaces and version control, it represents the future of computational biology. The interfaces will evolve: emerging tools like Claude Code Web promise to deliver agentic capabilities through browser-based environments, lowering the barrier to entry while maintaining full provenance tracking. As these tools mature, the combination of natural language interaction and automatic versioning will make reproducible AI-assisted analysis accessible to researchers regardless of their computational background.

Critically, all results produced with agentic AI tools require independent validation. In this study, we had the advantage of known expected outcomes—published results against which to benchmark our AI-assisted reanalysis. This “ground truth” allowed us to confirm that the AI-directed workflow produced biologically accurate results. However, for novel research where expected outcomes are unknown, researchers must exercise heightened scrutiny. AI agents can confidently produce plausible but incorrect interpretations, and without validation benchmarks, such errors may go undetected. We recommend orthogonal validation approaches: qRT-PCR confirmation of key findings, biological replication, functional studies, and cross-referencing with independent datasets. The provenance tracking enabled by agentic tools becomes especially valuable here—complete audit trails allow retrospective verification when questions arise about specific analytical decisions.

Why BRC-Analytics/Galaxy for AI-assisted analysis

The choice of workflow platform significantly impacts the feasibility of AI agent integration. Galaxy’s architecture offers distinct advantages over code-first systems like Nextflow for agentic AI workflows. Galaxy provides structured tool metadata through repositories like IUC (github.com/galaxyproject/tools-iuc), where each tool’s parameters, input/output types, and documentation are defined in machine-readable XML. This allows AI agents to query available tools, understand valid parameter options, and make informed decisions—capabilities that require parsing documentation or source code in Nextflow. Galaxy’s stateful API enables agents to inspect histories, monitor job status, and retrieve results through structured endpoints, whereas Nextflow requires log parsing and manual file path management. Perhaps most importantly, Galaxy’s integration with ACCESS-CI provides free, zero-configuration access to high-performance computing, eliminating the infrastructure barriers (container configuration, HPC authentication, resource allocation) that Nextflow imposes on users.

These architectural differences have practical implications for democratizing AI-assisted genomics. Galaxy’s web-based interface means users need only a browser, while AI agents handle the complexity of tool selection and parameter configuration through the API. Nextflow’s code-generation paradigm,

while flexible, requires users to review generated DSL2 scripts, configure execution environments, and debug failures—skills that remain barriers for bench biologists. As AI agents become integral to computational biology, platforms that provide structured metadata, stateful APIs, and accessible infrastructure will enable broader adoption than those requiring programming expertise to operate.

Future goals

We envision tighter integration between BRC-Analytics, Galaxy, and agentic AI systems. Currently, our workflow requires manual coordination: launching analyses through BRC-Analytics, managing data in Galaxy histories, and directing AI agents via API calls. Future development will embed AI agents directly within the Galaxy interface, enabling researchers to describe analyses in natural language while the system automatically selects appropriate workflows, configures parameters, and interprets results. This integration will transform BRC-Analytics from a workflow launcher into an intelligent research assistant that guides users through complex multi-omics analyses, suggests appropriate statistical approaches, and flags potential issues—all while maintaining the rigorous provenance tracking that Galaxy provides.

Materials and Methods

Literature Survey and Data Source Identification

To quantify *C. auris* sequencing data, we analyzed complete NCBI SRA database for taxonomy ID 498019 (*Candidozyma auris*) accessed December 3, 2025. SRA metadata (Cauris_SRA.csv) contained 27,201 total runs across 237 BioProjects. RNA-seq represents 812 runs (3.0%) and 64 BioProjects (27.0%), with WGS dominating run counts (26,201 runs, 96.3%) but representing 168 BioProjects (70.9%). Average runs per project: RNA-seq 12.7, WGS 156.0.

To characterize methodology across published RNA-seq studies, we linked all 64 RNA-seq BioProjects to associated publications. For each BioProject, we queried EuropePMC REST API (<https://www.ebi.ac.uk/europepmc/webservices/rest/>) for papers mentioning BioProject accession in full text, and NCBI E-utilities ([elink.fcgi](https://eutils.ncbi.nlm.nih.gov/entrez/eutils/efetch.fcgi)) for direct BioProject-to-PubMed links. This identified 21 papers linked to 20 of 64 BioProjects (31%); 44 BioProjects had no linked publications (unpublished or preprint). For papers with PMC IDs (17/20), we retrieved full-text XML and extracted reference genome information by pattern matching (GenBank/RefSeq accessions, strain names, clade designations) and RNA-seq tools (aligners, quantification tools, DE packages). Results in Supplementary Table 2.

For re-analysis validation, we selected Santana et al. (2023) *Science* (PRJNA904261) [6] and Wang et al. (2024) *Nature Communications* (PRJNA1086003) [7].

WGS Data Contributor Analysis

To characterize sources of *C. auris* WGS data, we analyzed the “Center Name” field from SRA metadata for all 26,201 WGS runs. Organization names were extracted and aggregated by run count and unique BioProjects. Abbreviated center names were expanded using geographic location metadata (geo_loc_name field) to disambiguate state-level public health laboratories (e.g., “MDH_CSL” mapped to Maryland via “USA:Mid-Atlantic” region; “NSPHL” mapped to Nevada via “USA:Nevada” location). Organizations were categorized into: US State/Local Public Health Laboratories, CDC, International Public Health agencies, Academic/Research institutions, and Other. Results presented in Supplementary Table 1.

Reference Genome and Annotation

All analyses used *Candidozyma auris* B8441 reference genome GCA_002759435.3 obtained from NCBI Assembly database. GTF annotation file contained 5,593 genes. This represents most recent annotation version at time of analysis and corresponds to assemblies used in BRC-Analytics platform.

RNA-seq Data Processing

Raw sequencing data (FASTQ files) for both BioProjects were obtained from NCBI SRA via BRC-Analytics platform. Standard pre-processing pipeline included: (1) Quality assessment using **FastQC**, (2) Adapter trimming and quality filtering using **fastp**, (3) Alignment to reference genome using **STAR** aligner, and (4) Gene-level quantification using **featureCounts**. All tools were executed through Galaxy platform (<https://usegalaxy.org>) using IWC workflows.

Counting Workflow

The BRC-Analytics platform executes Intergalactic Workflow Commission (IWC) workflows designed for reproducible RNA-seq analysis [10] (Figure 3). For paired-end data, the workflow begins with **fastp** for adapter removal and quality filtering, discarding reads shorter than 15 bp. Filtered reads are aligned to the reference genome using **STAR** with ENCODE-standard parameters, which simultaneously generates gene-level counts. Optionally, **featureCounts** provides an alternative counting method, while **StringTie** or **Cufflinks** compute normalized expression values (FPKM/TPM). Quality metrics from all steps are aggregated by **MultiQC** into a comprehensive report. The workflow also generates strand-specific coverage tracks (bigWig format) for genome browser visualization. All tools, versions, and parameters are locked within the workflow definition, ensuring identical results across executions.




Figure 3: IWC paired-end RNA-seq workflow. The pipeline processes FASTQ files through quality filtering (*fastp*), alignment (*STAR*), and quantification, with optional coverage track generation and QC aggregation via *MultiQC*.

Differential Expression Analysis

Gene count matrices from **featureCounts** were analyzed using **DESeq2** (v2.11.40.8+galaxy0) through Galaxy interface. For **Santana et al. dataset**: Samples organized into three collections (AR0382 n=2, AR0387 n=2, tnSWI1 n=2). Two pairwise comparisons performed: (1) AR0382 vs tnSWI1, (2) AR0382 vs AR0387. For **Wang et al. dataset**: Samples split into four collections by strain and condition (AR0382 *in vitro* n=3, AR0387 *in vitro* n=3, AR0382 *in vivo* n=3, AR0387 *in vivo* n=4). Two pairwise comparisons performed: AR0382 vs AR0387 in (1) *in vitro* and (2) *in vivo* conditions. **DESeq2** parameters: size factor normalization, Benjamini-Hochberg FDR correction, significance threshold FDR less than 0.01, fold change absolute value of log2FC greater than or equal to 1 for Wang dataset. Default parameters used for Santana dataset to match published analysis.

Gene Annotation Mapping

Published papers used older B8441 annotation versions with 6-digit gene ID suffixes while our analysis used GCA_002759435.3 with 5-digit suffixes. To reconcile gene identities, we developed LFC-based correlation mapping: for each gene in published DEG list, we identified the gene in our analysis with most similar log2-fold-change value. Mapping quality assessed using Pearson and Spearman correlation coefficients, direction agreement percentage, and mean LFC difference. Mapping scripts available at https://github.com/nekirut/claude-projects/tree/main/rnaseq/santana24_PRJNA904261/analysis and https://github.com/nekirut/claude-projects/tree/main/rnaseq/wang24_PRJNA1086003/analysis.

Galaxy Workflows and Reproducibility

All analyses performed on Galaxy Main server (<https://usegalaxy.org>). Galaxy histories containing complete analysis workflows, intermediate files, and final results are publicly accessible: - Santana et al.: <https://usegalaxy.org/u/cartman/h/prjna904261-final> - Wang et al. (Analysis): <https://usegalaxy.org/histories/view?id=bbd44e69cb8906b58b85fc3ebc05b72b> - Wang et al. (Final): <https://usegalaxy.org/histories/view?id=bbd44e69cb8906b59f131af7b542c1b1>

IWC workflows used are available at <https://iwc.galaxyproject.org> and are version-controlled in GitHub repository at <https://github.com/galaxyproject/iwc>. Workflow diagrams and analysis reports available in supplementary materials.

Statistical Analysis and Visualization

Validation statistics (Pearson correlation, Spearman correlation, direction agreement) calculated using custom Python scripts. Literature survey statistics and visualizations generated using `analyze_combined_data.py` and `visualize_combined.py` scripts. All analysis code and intermediate files available in project repositories.

Software Versions

- Galaxy platform: <https://usegalaxy.org> (accessed November-December 2024)
- DESeq2: 2.11.40.8+galaxy0
- STAR aligner, featureCounts, FastQC, fastp: versions specified in IWC workflows
- Python: 3.x (for validation scripts)
- Key Python packages: pandas, numpy, scipy, matplotlib, seaborn

References

1. Satoh K, Makimura K, Hasumi Y, Nishiyama Y, Uchida K, Yamaguchi H. *Candida auris* sp. Nov., a novel ascomycetous yeast isolated from the external ear canal of an inpatient in a Japanese hospital. Microbiology and Immunology [Internet]. 2009;53(1):41–4. Available from: <https://pubmed.ncbi.nlm.nih.gov/19161556/>
2. Centers for Disease Control and Prevention. Increasing threat of spread of antimicrobial-resistant fungus in healthcare facilities [Internet]. 2023. Available from: <https://www.cdc.gov/media/releases/2023/p0320-cauris.html>
3. *Candida auris*: A continuing threat. Microorganisms [Internet]. 2025; Available from: <https://pmc.ncbi.nlm.nih.gov/articles/PMC11946832/>

4. World Health Organization. *Candida auris* systematic review to inform the world health organization fungal priority pathogens list [Internet]. 2024. Available from: <https://pubmed.ncbi.nlm.nih.gov/38935900/>
5. National Institute of Allergy and Infectious Diseases. *Candida auris*—a mysterious and tenacious enemy [Internet]. 2024. Available from: <https://www.niaid.nih.gov/news-events/candida-auris-mysterious-and-tenacious-enemy>
6. Santana DJ, O’Meara TR, Romo JA, others. A *candida auris*-specific adhesin, scf1, governs surface association, colonization, and virulence. *Science* [Internet]. 2023;381(6665):1461–7. Available from: <https://pubmed.ncbi.nlm.nih.gov/37769084/>
7. Wang Y, Zou Y, Chen X, others. Cell surface glycan-lectin interactions modulate *candida auris* colonization and fungemia. *Nature Communications* [Internet]. 2024;15:6490. Available from: <https://pubmed.ncbi.nlm.nih.gov/38562758/>
8. Galaxy history: Santana et al. (2023) reanalysis [Internet]. 2024. Available from: <https://usegalaxy.org/u/cartman/h/prjna904261-final>
9. Galaxy history: Wang et al. (2024) reanalysis [Internet]. 2024. Available from: <https://usegalaxy.org/histories/view?id=bbd44e69cb8906b59f131af7b542c1b1>
10. Delisle L, Videm P, Erxleben-Eggenhofer A, Grüning B. RNA-seq-pe [Internet]. Zenodo; 2023. Available from: <https://doi.org/10.5281/zenodo.8354569>

Supplementary Materials

Supplementary Table 1: *C. auris* WGS data contributors by organization category and top sequencing centers.

Panel A: Summary by Organization Category

Category	Organizations	Runs	% of Total
US State/Local Public Health Labs	26	20,552	78.4%
CDC	2	2,626	10.0%
Academic/Research	46	1,365	5.2%
Other	41	1,345	5.1%
International Public Health	5	313	1.2%
TOTAL	120	26,201	100%

Panel B: Top 15 Contributing Organizations

Organization	Full Name	Runs	%
UPLH_ID	Utah Public Health Laboratory	4,447	17.0%
NVSPHL	Nevada State Public Health Laboratory	4,363	16.7%
CDC-NCEZID-MDB	CDC Mycotic Diseases Branch	2,406	9.2%
MDH_CSL	Maryland Dept of Health, Central Services Lab	2,309	8.8%
TXDSHS	Texas Dept of State Health Services	1,487	5.7%
MDHHS-GS	Michigan Dept of Health & Human Services	1,289	4.9%

Organization	Full Name	Runs	%
-	Wisconsin State Laboratory of Hygiene	1,211	4.6%
RIPHL	Rhode Island Public Health Laboratory	1,197	4.6%
NSPHL	Nevada State Public Health Laboratory	1,031	3.9%
-	Wadsworth Center (New York)	705	2.7%
-	Minnesota Dept of Health	688	2.6%
OCPHL_CA	Orange County Public Health Lab (California)	659	2.5%
-	Washington State Dept of Health	583	2.2%
UNLV NPM	Univ of Nevada Las Vegas, Pathogen Monitoring	443	1.7%
-	Fudan University	264	1.0%

US public health laboratories (state/local + CDC) account for 88.4% of all C. auris WGS data, reflecting outbreak surveillance priorities. Nevada appears twice (NVSPHL + NSPHL = 5,394 runs, 20.6%), indicating major outbreak focus.

Supplementary Table 2: RNA-seq methodology across 20 published *C. auris* BioProjects with linked publications (2018-2025).

BioProject	PMID	Authors	Year	Runs	Reference Genome	RNA-seq Tools
PRJNA445471	30559369	Muñoz JF et al.	2018	24	B8441, B11220, B11243	Bowtie2, TopHat2, RSEM, Trinity, edgeR
PRJNA477447	29997121	Kean R et al.	2018	22	B8441 (de novo)	Trinity, HISAT2, Kallisto, DESeq2
PRJNA682185	34630944	Zamith-Miranda D et al.	2021	36	B8441 (GCA_002759435.2)	DESeq2, edgeR
PRJNA682422	34180774	Lara-Aguilar V et al.	2021	6	B8441 (GCA_002759435.2)	FastQC, Trimmomatic, fastp, STAR, featureCounts, DESeq2
PRJNA735406	34354695	Zhou W et al.	2021	6	B11221 (Clades I-V)	Trimmomatic, HISAT2, Cufflinks, HTSeq, DESeq2
PRJNA801628	35473297	Biermann AR et al.	2022	24	B8441, B11221, B11243 (Clades I, III, IV)	HISAT2, featureCounts, edgeR
PRJNA830685	36445083	Narayanan A et al.	2022	16	B8441, CBS10913 (Clade II)	FastQC, fastp, BWA, Bowtie2, HTSeq, DESeq2
PRJNA788930	35652307	Shivarathri R et al.	2022	12	NS	RNA-seq
PRJNA792028	36913408	Bing J et al.	2023	15	GCA_002759435.2, GCF_002775015.1	HiSat2, StringTie, DESeq2, BWA
PRJNA904261	37769084	Santana DJ et al.	2023	6	B8441 (Clade I)	RNA-seq

BioProject	PMID	Authors	Year	Runs	Reference Genome	RNA-seq Tools
PRJNA1015296	38493178	Bing J et al.	2024	141	B8441 (GCA_002759435.2)	HiSat2, StringTie, DESeq2, BWA
PRJNA902676	38722168	Yang B et al.	2024	40	B11220, B11221 (Clades II, III)	Kallisto, DESeq2
PRJNA1036037	39480072	Li J et al.	2024	22	Clade IV	RNA-seq
PRJNA1086003	39455573	Wang TW et al.	2024	13	B8441 (Clade I)	HISAT2, STAR, DESeq2
PRJEB57846	39297640	Rhodes J et al.	2024	12	NS	WGS, RNA-seq
PRJNA1012821	40468551	Chauhan A et al.	2025	16	B8441, B11220 (CGD)	FastQC, fastp, Bowtie2, HTSeq, DESeq2
PRJNA1139166	40099908	Phan-Canh T et al.	2025	15	B8441 (GCA_002759435.2)	FastQC, fastp, cutadapt, STAR, featureCounts
PRJNA1208975	40530673	Yang G et al.	2025	9	Clade I	RNA-seq
PRJNA1232830	40066990	Chauhan M et al.	2025	6	Clade I	RNA-seq
PRJNA1291775	40863525	Vidal-Montiel A et al.	2025	6	GCA_003014415.1, GCA_034640365.1 (Clades III, IV)	FastQC, Trimmomatic, STAR, featureCounts, DESeq2

NS = Not specified in available text. Data extracted from PMC full text via EuropePMC and NCBI E-utilities APIs. 44 additional BioProjects had no linked publications.

Supplementary File 1: Galaxy workflow diagrams for re-analyses. *Source:* https://github.com/nekrut/claude-projects/tree/main/rnaseq/santana24_PRJNA904261/analysis and https://github.com/nekrut/claude-projects/tree/main/rnaseq/wang24_PRJNA1086003/analysis

Supplementary File 2: Gene mapping tables with LFC correlation values. *Source:* Same repositories as above

Supplementary File 3: Analysis reports with complete methodological details. *Source:* *ANALYSIS_REPORT.md* in each repository